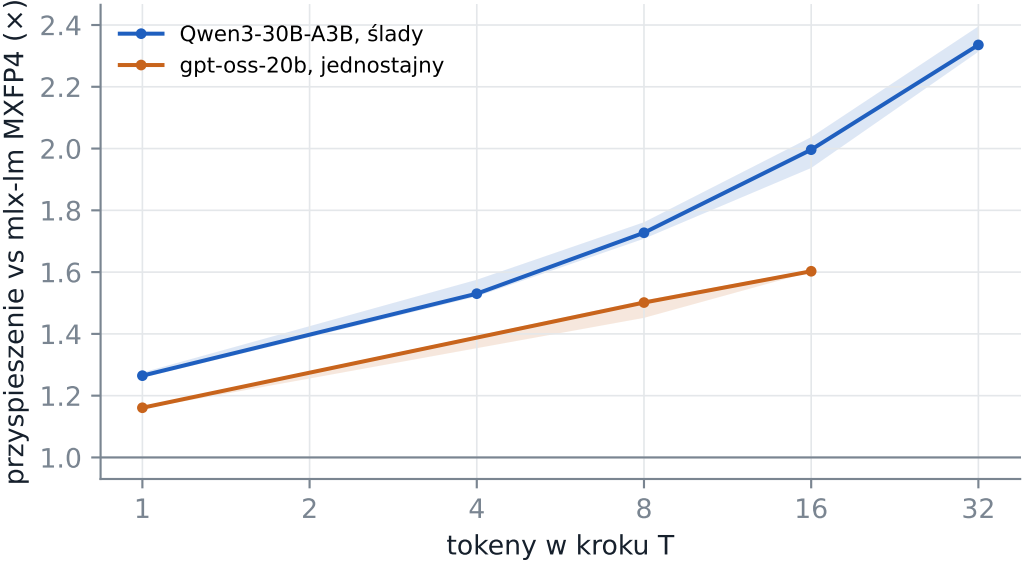
Qwen3-30B-A3B, ślady
gpt-oss-20b, jednostajny
przyspieszenie vs mlx-lm MXFP4 (×)
2.4
2.2
2.0
1.8
1.6
1.4
1.2
1.0
1
2
4
8
16
32
tokeny w kroku T