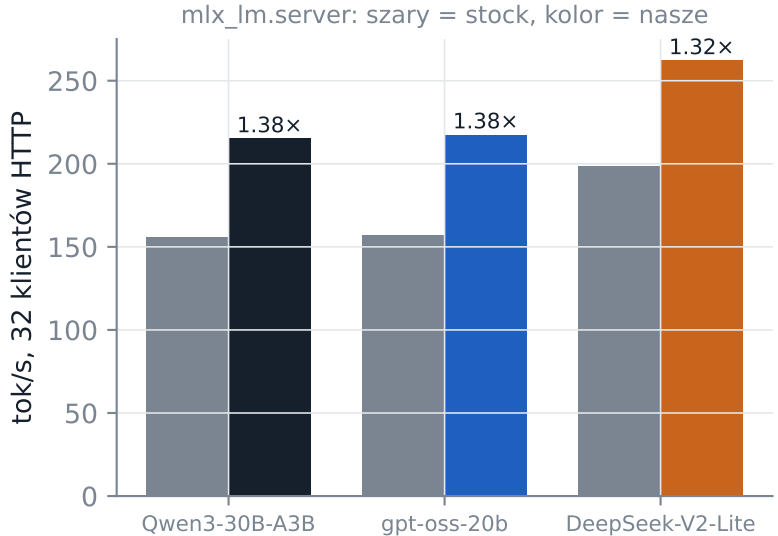
mlx_lm.server: szary = stock, kolor = nasze
tok/s, 32 klientów HTTP
0
50
100
150
200
250
1.38×
1.38×
1.32×
Qwen3-30B-A3B
gpt-oss-20b
DeepSeek-V2-Lite

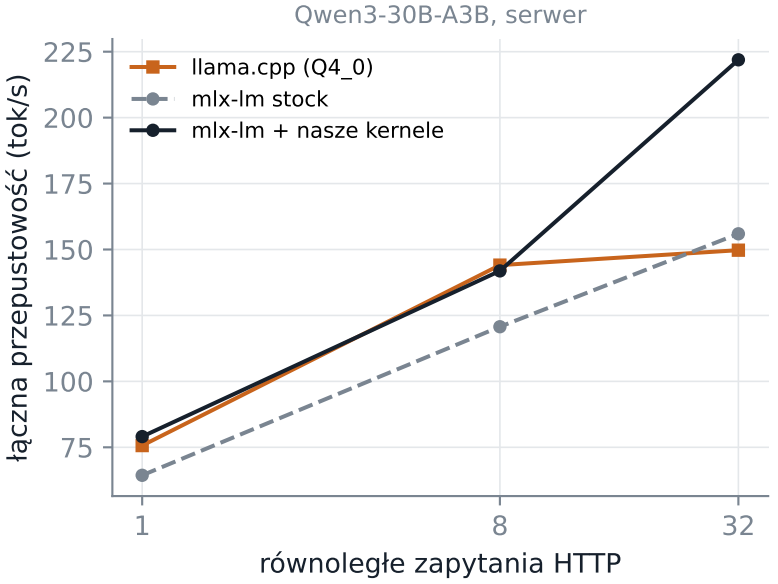
Qwen3-30B-A3B, serwer
łączna przepustowość (tok/s)
75
100
125
150
175
200
225
llama.cpp (Q4_0)
mlx-lm stock
mlx-lm + nasze kernele
1
8
32
równoległe zapytania HTTP